

Chapter 1

General Introduction

Sooner or later, every scientist or engineer runs into the problem of having to find the “best” solution some set of decisions can give. Usually, the variables governing such problems are related to one another in a rather complicated way, and finding the “best” combination of them can seem intractable. Normally, the relations between these decision variables can be translated into so-called *objective functions*, and the value of each of these functions can be interpreted as a measure of the quality that particular combination of the decision variables give to a particular aspect of the solution.

As such, N -dimensional optimization problems with m individual objective functions can be stated generally as

$$\text{find "min" } F(x) = \begin{bmatrix} f_1(x) \\ f_2(x) \\ \vdots \\ f_m(x) \end{bmatrix}$$

subject to

$$\begin{aligned} G(x) &< 0 \\ H(x) &= 0 \\ lb &\leq x \leq ub, \end{aligned}$$

where

$$x = [x_1, x_2, \dots, x_N],$$

and $f_1(x)$ through $f_m(x)$ indicate the individual objective functions. In almost all practical cases, finding good solutions is not a problem, but finding the *best* solution is much more difficult. Moreover, for most problems encountered in practice, the objective functions $f_1(x)$ through $f_m(x)$ are highly nonlinear,

non-smooth, non-differentiable, or have no way of determining initial estimates close to the global optimum, or a combination of all of these factors.

The only practical solution to tackle such problems is to use a so-called *meta-heuristic optimizer*, which uses a “population” of trial solutions, and applies certain probabilistic rules to generate a new population which converges to the global minimum of the objective function with high probability. Over the years, many such algorithms have been developed, of which the Genetic Algorithm (GA), Differential Evolution (DE), Particle Swarm Optimization (PSO) and Adaptive Simulated Annealing (ASA) received more attention.

Although meta-heuristic optimization algorithms usually require quite many function evaluations, their popularity grew enormously due to their ability to find global optima even in extremely difficult problems and with relatively low population sizes. Moreover, their elegance and simplicity appealed to many people – a thorough understanding of other optimization algorithms that existed at the time (which were much harder to implement, and did not always find the *global* optimum), was no longer required, and problems could be optimized much faster (and usually much better) than was previously possible.

Chapter 2

Single-objective Optimization

Originally, the aforementioned meta-heuristic algorithms were intended for problems with $m = 1$ (referred to as single-objective optimization); their aim is to find the global minimum of a *single* objective function. Many problems can indeed be stated as a single objective problem:

$$\begin{aligned} & \text{find } \min F(x) \\ & \text{subject to} \end{aligned}$$

$$\begin{aligned} G(x) &< 0 \\ H(x) &= 0 \\ lb &< x < ub \end{aligned}$$

and $x = [x_1, x_2, \dots, x_N]$ as before. Meta-heuristic algorithms first create a population of randomly generated solutions,

$$pop = \begin{bmatrix} x_{11} & x_{21} & x_{31} & \dots & x_{N1} \\ x_{12} & x_{22} & x_{32} & \dots & x_{N2} \\ & & \vdots & & \\ x_{1P} & x_{2P} & x_{3P} & \dots & x_{NP} \end{bmatrix}$$

where each x_{ij} is taken within the preset boundaries $[lb]$ and $[ub]$ (the constraints $G(x) < 0$ and $H(x) = 0$ are then usually added to $F(x)$ in the form of *penalty functions*). The objective function $F(x)$ is evaluated for each member in this population, and a new population is created based on the function values of the initial population, and a certain degree of randomness. The four aforementioned algorithms do this as follows:

2.1 GA (based on natural evolution)

1. Select two individuals that function as parents, say individuals 2 and 8.
2. split the parents in two, at some random location CR :

$$\begin{aligned}parent1 &= [x_{18}, x_{28}, \dots, x_{CR8}, \dots, x_{N8}] \\parent2 &= [x_{12}, x_{22}, \dots, x_{CR2}, \dots, x_{N2}]\end{aligned}$$

3. Let the parents *crossover* at the point R (with a certain probability p_{cross}) to create two children:

$$\begin{aligned}child1 &= [x_{18}, x_{28}, \dots, x_{CR2}, \dots, x_{N2}] \\child2 &= [x_{12}, x_{22}, \dots, x_{CR8}, \dots, x_{N8}]\end{aligned}$$

4. Do this until P children have been created.
5. *Mutate* the children, with a certain (small) probability p_{mutate} . This selects a few random indices (M) in ALL children, and replaces the associated values with random other values (randomly selected from the interval $[lb, ub]$):

$$\begin{aligned}child1 &= [x_{18}, x_{28}, \dots, x_{CR2}, \dots, x_{M2}, \dots, x_{N2}] \\child2 &= [x_{12}, x_{22}, \dots, x_{CR8}, \dots, x_{N8}]\end{aligned}$$

6. Evaluate the objective function for all the children. If a child is found to have a “better” function value than either of its parents, it will become part of the new population. Otherwise, the better of the two parents is inserted into the new population.

Steps 1-6 are repeated until “convergence”.

The original GA used a *binary* representation of the population, i.e., each individual is represented by *bits* instead of real numbers. Crossover and mutation are also carried out bit wise, that is

```

parent1 = [10110011010010...00110110001001]
parent2 = [11001110110011...00110011111001]
crossover →
child1 = [10110011010010...00110011111001]
child2 = [11001110110011...00110110001001]
mutation →
child1 = [10110011010010...00100011111001]
child2 = [11011110110011...00110110001000]

```

Whether to use binary representation or real numbers usually depends on the problem. For lower dimensionality (N is small) it is usually more efficient to use binary representation, and when the population size is enormous it is more efficient to use real numbers to avoid the costly conversion to binary and back, etc. But these “rules-of-thumb” usually need to be tested for each new problem.

2.2 DE (based on globalized pseudo-derivatives)

1. Randomly select three individuals from the population, say 3, 7 and 15. These individuals will function as the *base vector* and *differentiation vectors*, respectively.
2. The i^{th} individual is created according to the rule

```

if rnd < Cr
    ind = pop(3, :) + F(pop(7, :) - pop(15, :))
else
    ind = pop(i, :)
end

```

where **rnd** is a random number, Cr is the *crossover probability*, and F the *constant of differentiation*, usually a random number in $[-1, 1]$.

3. Do this until P new individuals have been created.
4. Evaluate the objective function for all these new individuals. If a new individuals is found to have a “better” function value than its spawning solution, it will become part of the new population. Otherwise, the original vector is inserted.

Taking the *difference* between the two differentiation vectors is very much like taking the derivative. But as the two differentiation vectors are usually quite far apart (certainly not infinitesimally far), this “derivative” is more a

global measure of how much the objective function changes *on average* over that interval. The derivative is computed at each iteration between two new, *randomly* selected vectors, so on average, the solutions will tend to go to where the average slope is zero, and the function globally minimal. Sometimes this operation is called the *global pseudo-derivative*, and it is the key to the power of the DE algorithm.

Numerous variations on this basic algorithm exist. However, they normally improve DE's performance only marginally, and as such, they will not be mentioned here.

2.3 ASA (based on the laws of thermodynamics)

1. Randomly perturb every individual in the population. The so-called Boltzmann generating scheme accomplishes this:

$$ind = ind + \text{sqrt}(T) * \text{randn}(1, \text{dimensions}),$$

with `randn()` random numbers from the standard normal distribution, and T the current temperature.

2. Evaluate the objective function for all new individuals.
3. Accept or reject new individuals into the next population. If the value of the objective function is lower than before the perturbation, always accept it. If it is higher, accept it according to the probabilistic rule

$$\text{accept if } rnd < \exp \frac{E_0 - E_p}{T},$$

where $E_0 - E_p$ is the difference in objective function values before (E_0) and after (E_p) the perturbation, and T is the current temperature.

4. At every new iteration, the temperature is first decreased according to a cooling schedule. Usually, this cooling schedule has the form

$$T_{\text{new}} = c \cdot T_{\text{old}},$$

where $0 < c < 1$ is a constant. This form will decrease the temperature logarithmically, just as it would in physical system undergoing cooling¹.

¹The algorithm described above is how it is implemented in GODLIKE. Do note that the algorithm described above is by no means *adaptive* – I just gave it that title as a placeholder for future work to be carried out. Adaptive SA means that the algorithm automatically adjusts the cooling schedule and the Boltzmann constant to optimize the quality of the converged solution. But as of yet, that is not yet implemented.

Traditionally, for this method, individual trial solutions are called “atoms” or “particles”, to reflect the method’s underlying philosophy – as the temperature drops, the atoms literally “freeze” into low-energy states (low function values). But before they freeze, they have the ability to move to *higher* energy states, with a certain probability (step 3). This is what makes ASA also a global optimizer, in the sense that it is not “greedy” as to only accept lower function values, but also explores regions behind high-energy barriers.

Originally, Simulated Annealing was built around a single solution (the initial condition). However, this method is easily rewritten into a population-based method (just use N randomly generated initial conditions).

2.4 PSO (based on swarm intelligence)

1. Aside from a population of randomly generated initial trial solutions, also initialize for every individual a *velocity* V in an arbitrary direction of the same dimensionality as the problem. Also create a small social network for every individual, by assigning a number of “neighbors” or “friends” to each individual. These are just a number of other individuals associated to one individual, that influence the individual.
2. New individuals are generated every iteration simply by adding the step associated with the current velocities for each individual, e.g.,

$$pop_{i+1} = pop_i + V_i$$

3. Evaluate the objective function for all new individuals thus created.
4. Keep track of three values per individual: *lbest*, *nbest* and *gbest*. The value *lbest* is the *local best* function value, that is, the best function value ever encountered by each individual, and the associated location where it encountered it. The value *nbest* is the *neighbor best*, or the best function value (and its location) encountered by each of an individual’s neighbors. Finally, *gbest* is the *global best*, that is, the best function value (and location) ever encountered by *all* individuals.
5. update the velocity according to the rule

$$\begin{aligned} V_{i+1} = & \omega V_i + \\ & + \text{rnd}_1 \cdot \eta_1 \cdot (ind_i - nbest) \\ & + \text{rnd}_2 \cdot \eta_2 \cdot (ind_i - gbest) \\ & + \text{rnd}_3 \cdot \eta_3 \cdot (ind_i - lbest), \end{aligned}$$

where ω is the inertia constant, η_1 is the *social learning factor*, η_2 is the *cooperative factor*, η_3 is the *cognitive learning factor*, and rnd_{1-3} are three random numbers from $[0, 1]$.

The last step is the crux of the algorithm. Updating velocities in this fashion will steer every particle into a direction that was found to be good by its neighbors (social learning), a direction found to be good by all individuals combined (cooperative), and a direction that each individual found to be good in the past (nostalgia). This gives the particles (the traditional name for individuals) a type of behavior reminiscent of a swarm of insects around a good food reserve – most swarm around it, having a feeding frenzy (local optimization), while others remain swarming in a relatively large area around it (localized global search), and sometimes there are the true explorers going to completely new areas (global search).

Chapter 3

Multi-objective Optimization

The power and popularity of these single-objective optimization algorithms encouraged many to re-state their optimization problems with multiple objectives, to have only *one* objective function, usually something like

$$K(x) = f_1(x) + af_2(x) + bf_3(x) + \dots + Zf_M(x)$$

or

$$K(x) = f_1^2(x) + f_2^2(x) + f_3^2(x) + \dots + f_M^2(x)$$

or something similar. However, when two different such rules are used to convert the multi-objective problem into a single-objective one, they usually also give *different* solutions. Creating a single-objective problem in the way described above needlessly introduces extreme sensitivity to the problem specifics, which is far from desirable. More importantly, the “optimum” this process generates may not be desirable at all. This is best illustrated by an example.

Consider the optimization of the trajectory of a spacecraft to the planet Mars. Usually, for scientific missions, the main goal is to get the largest amount of mass on Mars (equal to minimizing $-f_{mass}(x)$). But when considering *manned* missions, the *time* the spacecraft takes to get there is also very important, as essential consumable resources (oxygen, food, water, ...) are limited, and long exposures to space far from Earth’s protective magnetosphere will cause all sorts of illnesses in the crew. Thus, finding the minimum time is also highly desirable (equal to minimizing $f_{time}(x)$).

In theory it is quite possible to go to Mars in only one week – just bring an enormous rocket (large accelerations aside). However, to go there in only one week requires *vast* amounts of propellant, so that almost *nothing* will be left

in Mars orbit, let alone for the return trip. On the other hand, going to Mars without using *any* propellant is also possible ((re)entry and launch aside); just use the “interplanetary super highway”, made possible by the combined gravitational effects of all the planets, and you have a free ride to Mars. However, even well-selected trajectories along this superhighway take longer than 5 years to get to Mars, which makes it near-impossible with today’s engineering to have any survivors (let alone volunteers for such a mission).

Optimizing the *sum* of both these objectives will nearly always benefit the short-time criterion more than the high-end-mass criterion, or vice versa, in other words, introduce a bias towards one of the objectives. Therefore, the particular choice of summation is highly problem dependent, and it should exhaustively be experimented with before any valuable results can be obtained. This process is extremely tedious, and should be re-done for *every* new optimization problem.

For such problems, it is actually most desirable to get the best *compromise* between the two objectives, in stead of the optimum of the most un-biased sum of both. Moreover, as optimization goes in large-scale projects, it is also desirable to have a *set* of different good compromises, so that these costly and lengthy optimizations do not have to be done all over again in case something in the project changes.

The set of best compromises are usually given in the form of the associated *Pareto front*. An easy example is the best compromise between the functions $\sin(\mathbf{x})$ and $\cos(\mathbf{x})$. The associated Pareto front is shown in Figure 3.1, and the (rather obvious) location of the very best compromise is shown in Figure 3.2 (Note that the Pareto front is in function value space (function values are plotted against each other)).

There are several algorithms that find the complete Pareto front in multi-objective optimization problems. The most popular, easiest to implement and most efficient one known, still is the Non-dominated Sorting Genetic Algorithm II (NSGA-II). This algorithm sorts the current population according to the amount of solutions that *dominate* each other individual, Dominance of one individual x_i over another y_i , denoted as $x_i \prec y_i$, is defined as

$$\begin{aligned} x_i \prec y_i \quad & \text{if} \quad f_j(x_i) \leq f_j(y_i) \quad \text{for all functions } j, \\ & \text{and} \quad f_j(x_i) < f_j(y_i) \quad \text{for at least one function } j. \end{aligned}$$

The NSGA-II algorithm iterates the following steps until all solutions are non-dominated:

1. Create an offspring population Q from the parent population P with the usual crossover and mutation operators from a GA.
2. Count the number of solutions y_i that dominate the current solution x_i .

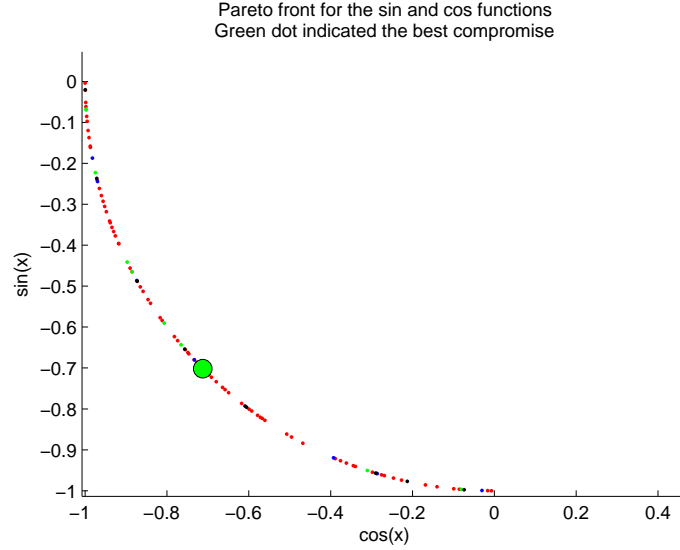


Figure 3.1: All good compromises for the minimum of the *two* objective functions $\sin(x)$ and $\cos(x)$, in function value space. Note that the green dot is closest to the origin, and thus indicates the “most efficient” compromise.

Do this for *all* individuals from *both* the parent population P and the offspring population Q .

3. Some solutions will be found to have *zero* other solutions dominate them. They are *non-dominated*, and thus part of the Pareto front of the current populations. The solutions that have only one other solution dominate them, would have been part of the Pareto front *if* the members forming the true Pareto front would not have been present. Those that have two solutions dominate them would have formed the Pareto front if *those* solutions would also not be present, etc. Thus, the level of domination is indicative of the quality of that solution.
4. Next, the *crowding distances* are computed. These are the *average* distances between one solution and its surrounding solutions in the function-value space.
5. Create a new population R , which contains individuals from the previous two populations P and Q , sorted by their level of dominance. That is, first insert all Pareto members in R , then those that have only one dominating solution, etc. Keep inserting individuals until R is the same size as P and Q .
6. Create a subset P_{i+1} from R by a binary *tournament selection*. This se-




Figure 3.2: Best compromise for the minimum of the *two* objective functions $\sin(x)$ and $\cos(x)$, in decision variable space.

lection takes two random individuals from R , a_R and b_R , and lets them *compete* using their domination level and crowding distances as competitive factors. The “winning” individual is the one that satisfies $a_R \prec_d b_R$, defined as

$$\begin{aligned}
 a_R \prec_d b_R \quad & \text{if} \quad \text{rank}(a) < \text{rank}(b) \\
 & \text{or} \quad (\text{rank}(a) = \text{rank}(b) \\
 & \text{and} \quad \text{crowding_distance}(a) > \text{crowding_distance}(b))
 \end{aligned}$$

where $\text{rank}(\ell)$ indicates the rank, or domination level, of the individual ℓ . This process is repeated until the subset S is full. Usually, the size of P_{i+1} is taken to be half that of Q and R .

7. Create a new offspring population Q_{i+1} , equal in size as the original P , Q and R , using crossover and mutation from a GA, using members from the subset P_{i+1} as parents.

After the initialization step 1, steps 2 through 7 are repeated until *all* individuals are non-dominated. The crowding distances in steps 4 and 6 are used to keep the *spread* in the solutions along the true Pareto front more or less homogeneous – when these steps are not included, the solutions tend to cluster together to the easiest-to-find compromise between the objective functions.

The greatest advantage of NSGA-II is that the *entire* population will simply converge to the *true* Pareto front, so that the number of desired solutions can easily be controlled by choosing a different population size. A slight drawback it has compared to other algorithms of this sort is the computational complexity of the computation of the number of non-dominated solutions; it is of $O(M(2N)^2)$, where M is the number of objectives and N is the population size. With careful bookkeeping, this can be reduced to $O(MN^2)$, but still it tends to be a problem for very large population sizes. However, with today's standards in computation power this really poses only minor problems.

Note that the genetic operators used to create Q or Q_{i+1} are completely separate from the other parts of the algorithm, so Q and Q_{i+1} can essentially be generated with *any* of the aforementioned meta-heuristic optimizers. This fact will be used later on.

Chapter 4

Problems with Meta-heuristic Algorithms

Despite their popularity and general applicability, there are many problems associated with meta-heuristic optimizers. One very serious problem is *premature convergence* – the population converges to a point that is only a *local* minimizer of the function. Also, the NSGA-II algorithm used for multi-objective problems might return fully non-dominated solutions, but completely miss the problem’s Pareto front; non-dominance is not a guarantee for convergence to the Pareto front. This is particularly true for small population sizes. For small populations, the probability that no solution is dominated by another, while still not being even close to the Pareto front, is quite large. This necessitates using large population sizes, thus requiring many function evaluations.

The GA, DE, PSO and ASA algorithm have various operators that try to prevent this problem. In GA, it is the mutation operator that sometimes generates a solution very far from the other population members, increasing its robustness. In PSO, it is the fact that sometimes solutions get assigned very large changes in their velocities, especially when being very far removed from the attractors. However, as found in the literature, premature convergence still frequently occurs and necessitates several re-runs to make sure the global optimum is found.

A fact that can also not be ignored is that people tend to “get used” to the simplicity of such algorithms. It often becomes their algorithm of choice for *all* optimization problems they encounter. Quite often, obvious caveats, simplifications and weaknesses in the problem that can be exploited by *much* more powerful and accurate algorithms, are then “overlooked” or sometimes even ignored.

It should be stressed that meta-heuristic algorithms do *not* aim to find the global minimum very accurately, nor do they aim to be *efficient* in terms of function

evaluations; they only aim to find a good *approximation* to the problem’s global minimum or Pareto front, with high probability. For many problems, function evaluations are *very* expensive; personally I frequently encounter functions that take several *minutes* to evaluate per trial, even on a 32-node quadcore cluster (128 processing units). In such cases it is much more fruitful (and efficient) to analyze that problem to bits, and make some reasonable assumptions that simplify the problem significantly, rather than go at it “blindly” and just use a GA to solve it.

Despite these problems, they are still very useful, if only to get a good initial approximation to the minimum, or a first idea about the general shape of the Pareto front. They can also indicate other promising regions in the search space, or function as a simple test algorithm to find potential problems with the objective functions.

Chapter 5

GODLIKE Algorithm

The GODLIKE algorithm was written as an attempt to improve the robustness of the meta-heuristic algorithms, and to do away with the need to fine-tune the algorithm of your choice for each optimization problem. It was also written to serve as a general “umbrella” function; to be able to tackle both single and multi objective problems with a single function and in a uniform fashion, and easily include more and different population based methods.

GODLIKE stands for **G**lobal **O**ptimum **D**etermination by **L**inking and **I**nterchanging **K**indred **E**valuators, and this is exactly what it does. It uses all four aforementioned algorithms simultaneously (Linking), and after convergence of either of them, or exceeding certain predefined limits, it takes random members from each population and inserts them into random other populations (Interchanging) before continuing the optimization.

By using multiple optimizers simultaneously, it is essentially equal to performing four (or more) consecutive optimizations all at once, which already improves the chances of finding the global optimum; The weaknesses associated with each algorithm are negated by the strengths of another, while the strengths of all algorithms simply add up.

The *interchange*-operator indeed destroys part of the convergence properties of either of the algorithms it uses, but that is exactly the intention – the convergence one of the algorithms is experiencing might be to a local optimum, while the others might be converging to the global solution, or other local minima. By interchanging individuals between populations, GODLIKE introduces *immigrants* into the populations that can provide alternative good solutions to the ones already being explored by one of the algorithms. These immigrants can steer the population into other, unexplored areas of the search space, increasing the chances of locating the global minimum. By keeping the populations separate, also the principle of *isolation* is exploited automatically – portions of the search space will be thoroughly explored by one of the populations, while not

affecting the other populations.

The interchange operator is extremely useful for multi-objective problems; when one population is completely non-dominated, interchanging individuals between populations will usually result in a *dominated* population, which continues the search for the Pareto front, instead of reporting convergence.

In conclusion, GODLIKE does not aim to make either of the algorithms more efficient in terms of function evaluations, (rather, it tends to require *more* function evaluations). However, the *robustness* was aimed for, and until now, my simple experiments have indeed shown that hard-to-find global optima that could almost never be found by either GA, DE, ASA or PSO individually, *could* be found by but their *combined* efforts in GODLIKE.

5.1 GODLIKE in Detail

GODLIKE was written primarily with all of the above in mind, but also partly for me personally to *finally* learn objective-oriented programming in MATLAB. As such, the files `pop_single.m` and `pop_multi.m` are indeed `classdef`-class definitions. A slight drawback of this is that only users who own MATLAB 2008b (or later) can use it, but I think it is not hard (only time-consuming) to re-write these files to pure functions.

GODLIKE requires four files: `GODLIKE.m`, `pop_multi.m`, `pop_single.m` and `set_options.m`.

5.1.1 GODLIKE.m

This is of course the main function. All required operations are carried out here. The operation of GODLIKE.M is kept simple, readable and understandable by generously using nested-functions. All basic operations are performed in the first three cells:

%%Initialize Here, the user-input is checked thoroughly. Also, the user-provided objective function(s) are tested and used to determine whether single or multi-objective is desired. Also, the user-input is reshaped and reformatted into a form that is assumed in all classes and functions. During this last process, also default options and values are assigned should they be empty or omitted.

%%GODLIKE loop The main loop that executes the optimization. In this loop, the amount of individuals per optimizer, and the number of iterations that is to be carried out by each optimizer, is selected by randomly “breaking up” the user-selected (or default) values. For example, if 100 individuals are to be used in the algorithms GA, DE, and PSO, the algorithms get assigned for instance [45, 13, 42] or [6, 22, 72] individuals.

Note that these numbers are chosen differently at every iteration.

Next, using these values, the number of desired populations are created (objects of type `pop_single.m` or `pop_multi.m` are instantiated). Then, GODLIKE performs the randomly selected number of iterations in each algorithm, keeping track of the convergence of either of them. For single objective optimization, convergence is said to have been achieved if the decrease in the global minimum found so far is less than `[options.MinDescent(1)]` for at least `[options.MinDescent(2) * options.MinDescentMultiplier]` iterations (see `set_options` below). For multi-objective optimization, convergence of an algorithm occurs simply when all members of that population are non-dominated.

After all algorithms have been run for the said amount of iterations, convergence of the GODLIKE loop is checked in a similar manner: for single-objective optimization, if the amount of GODLIKE iterations is larger than `[options.MinIters]` and the decrease in the global optimum is less than `[options.MinDescent(1)]` for at least `[options.MinDescent(2)]` GODLIKE iterations, the GODLIKE loop is terminated. For multi-objective problems, at least `[options.MinIters]` GODLIKE iterations will be performed, and if after said amount of iterations all solutions are non-dominated, the GODLIKE loop is terminated. Note that the GODLIKE loop is also terminated when more than `[options.MaxIters]` iterations have been performed, or when more than `[options.MaxFunEvals]` function evaluations have been executed at any point in the loop.

Every next GODLIKE iteration again randomly selects the number of individuals and iterations per algorithm to be carried out. Only this time, the existing populations are first *shuffled*; that is, the *interchange*-operator is applied.

%%output values Simply assigns the desired output variables. These are updated every iteration, but only properly formatted and nicely cut into pieces by the operations in this cell, before they are returned to the user.

Note that the operation manual (how to actually use GODLIKE in MATLAB) is included in a separate PDF-file (`manual.pdf`). The nested functions are where the actual work is carried out. Note also the nested function `display_progress`: this is called only when `[options.display]` is set to 'plot' or 'on'.

5.1.2 `pop_single.m`

A SubClass of the `handle` class, this is the file where all the actual optimizations are carried out. It constructs a "population" called `pop`, which has properties

algorithm the optimization algorithm used (either 'GA', 'PSO', 'DE' or 'ASA').

funfcn The objective function.

individuals All members of the population.

fitnesses their corresponding objective function values

size population size (number of individuals).

lb lower bounds,

ub upper bounds. Note that both the lower and upper bounds are replicated and resized upon initialization, to conform to the size $[\text{popsize} \times \text{dimensions}]$. This does away with the need to constantly replicate them for operations like **check_bounds**, or re-initializing individuals as is done in some of the algorithms.

dimensions The dimensions of the problem.

funevals Total number of function evaluations made

iterations Total number of iterations so far performed

options A copy of the options structure (see **set_options**)

pop_info A structure to store intermediate data and pass it from method to method. For single-objective optimization, it contains the following fields:

- **parent_population**
- **offspring_population**
- **function_values_parent**
- **function_values_offspring**

Note that the fields **parent_population** and **function_values_parent** are only here for completeness and consistent programming, their contents is completely the same as the class properties **individuals** and **fitnesses**.

These properties are all assigned by the constructor (when properly called). The constructor also creates the initial population of randomly generated individuals within the given bounds, and evaluates the function for these initial individuals. Note that the constructor does not perform any elaborate checks on the given input; it relies on the fact that this has already been done in **GODLIKE.m**.

On a population **pop**, the following methods can be applied:

iterate (Public) Perform one single-objective iteration. Aside from the constructor, this is actually the only method directly accessed in **GODLIKE.m**; all other methods are only used *within* **pop_single.m**.

create_offspring (Hidden) Creates offspring from the parent population, using the pre-set algorithm. These offspring are inserted into the `pop_info` structure. Note that this contains only the *first half* of all the optimization algorithms.

evaluate_function (Hidden) Proper evaluation of the objective function. Function can be evaluated in two distinct ways, each of these requiring implementation in `pop_single`. The correct one has been determined in GODLIKE, which is passed to `pop_single` in the background via the `options` structure.

replace_parents (Hidden) Selective replacement of the parent population; the exact replacement procedure depends on the selected algorithm. Note that the *second half* of the algorithms is carried out here.

check_bounds (Hidden) Checks whether offspring is generated within the given bounds `lb` and `ub`. Also, for PSO specifically, the bounds on the new velocities are checked.

initialize_algorithms (Hidden) PSO and ASA need some additional initialization (temperature, velocities, ...). With future expansions of GODLIKE in mind, I thought I make a separate method for this.

5.1.3 pop_multi.m

A SubClass of `pop_single.m`. Inherits all the methods defined therein, and is constructed in exactly the same way. It adds only one property `num_objectives`, and the two methods `non_dominated_sort` and `tournament_selection`.

Note that the methods `evaluate_function`, `initialize_algorithms` and `iterate` are overloaded in `pop_multi.m`. This is required because multi-objective functions need to be evaluated differently (there may be multiple functions, or a single function with two-dimensional output), one iteration now must call `non_dominated_sort` and `tournament_selection` instead of `replace_parents`, and the PSO-algorithm must be initialized and used differently (see below).

5.2 PSO in Multi-objective Optimization

Using the ASA and DE optimizers (in stead of the usual GA) to create offspring populations for the NSGA-II method is quite straightforward. However, the PSO algorithm is more demanding; it needs to have some values for *lbest*, *gbest* and *nbest*, which are not defined in an obvious way for multi-objective optimization. There are several papers in the literature that deal with this very problem. Some find reasonable results, but most perform quite poorly compared to similar algorithms. The best one I could find relied on a completely different algorithm (so *not* NSGA-II), so for the time being, I chose to use another, worse one, that *did* fit in GODLIKE's NSGA-II context.

But, it is particularly in this context that PSO seems rather unsuited for the task. To remedy the problem with *lbest*, *gbest* and *nbest*, the following criteria seemed to work best:

- only update *lbest* if the new individual dominates the previous *lbest*, and if it is part of the current population's Pareto front (*rank* = 0)
- only update *gbest* with one of the members of the current Pareto front, that has a larger crowding distance than the previous *gbest*.
- For every iteration, find the `options.num_neighbors` individuals that are closest in function-value space to the current individual. These form the individual's *new neighbors*. The best of these neighbors, *nbest*, dominates all other neighbors and has the largest crowding distance.

Despite these remedies, PSO remains the least powerful algorithm for multi-objective optimization; it doesn't really seem to be able to achieve convergence.

5.3 Known Problems and Issues

As mentioned above, using the PSO algorithm on multi-objective problems should be done sparingly, or at least in conjunction with another algorithm. Initially it seems to push the solutions towards the Pareto front quite fast, but I think this is more due to the non-dominated sorting and tournament selection. It is very hard to get convergence with PSO, but in conjunction with DE it does work reasonable.

The 'A' in (A)SA is really not deserved. Basically I just used the simulated annealing written by Joachim Vandekerckhove (also on the FEX, file ID#10548) and rewrote it to be suited for populations. Adapting the control parameters at each iteration, which is the true power of ASA, is not included now. As such it is probably the weakest algorithm for single-objective optimization, and the second-worst in multi-objective problems. However, I found it to still be useful sometimes, as it generates many solutions in *all* regions of low function value.

For very large population sizes, the crossover operator in GA seems to take up very much computation time. I have no idea why that is, so if you do know why, please let me know. Also, the to-and-fro conversion between binary and real representations is relatively costly. While this is probably not a problem (the effect only becomes noticeable for population sizes larger than ~ 4000), it's still something I'm baffled about. If anyone knows a better way to do it, please let me know.

5.4 Future Work

Currently, GODLIKE only accepts objective functions, be it one or many. In the near future, I think it is most important to include the possibility to also pass it constraint functions. In the current form, penalty-function methods must be used to incorporate constraint functions, but this is certainly not the best way to do it.

Naturally, I need to do some more research on using PSO in multi-objective problems. It was found to be quite promising in the literature despite the implementation difficulties, but the current implementation in GODLIKE does not really solve the problems satisfactorily.